

LAB: Working with Hadoop and MapReduce

Outline

HDFS

- creating folders
- copying files
- ...

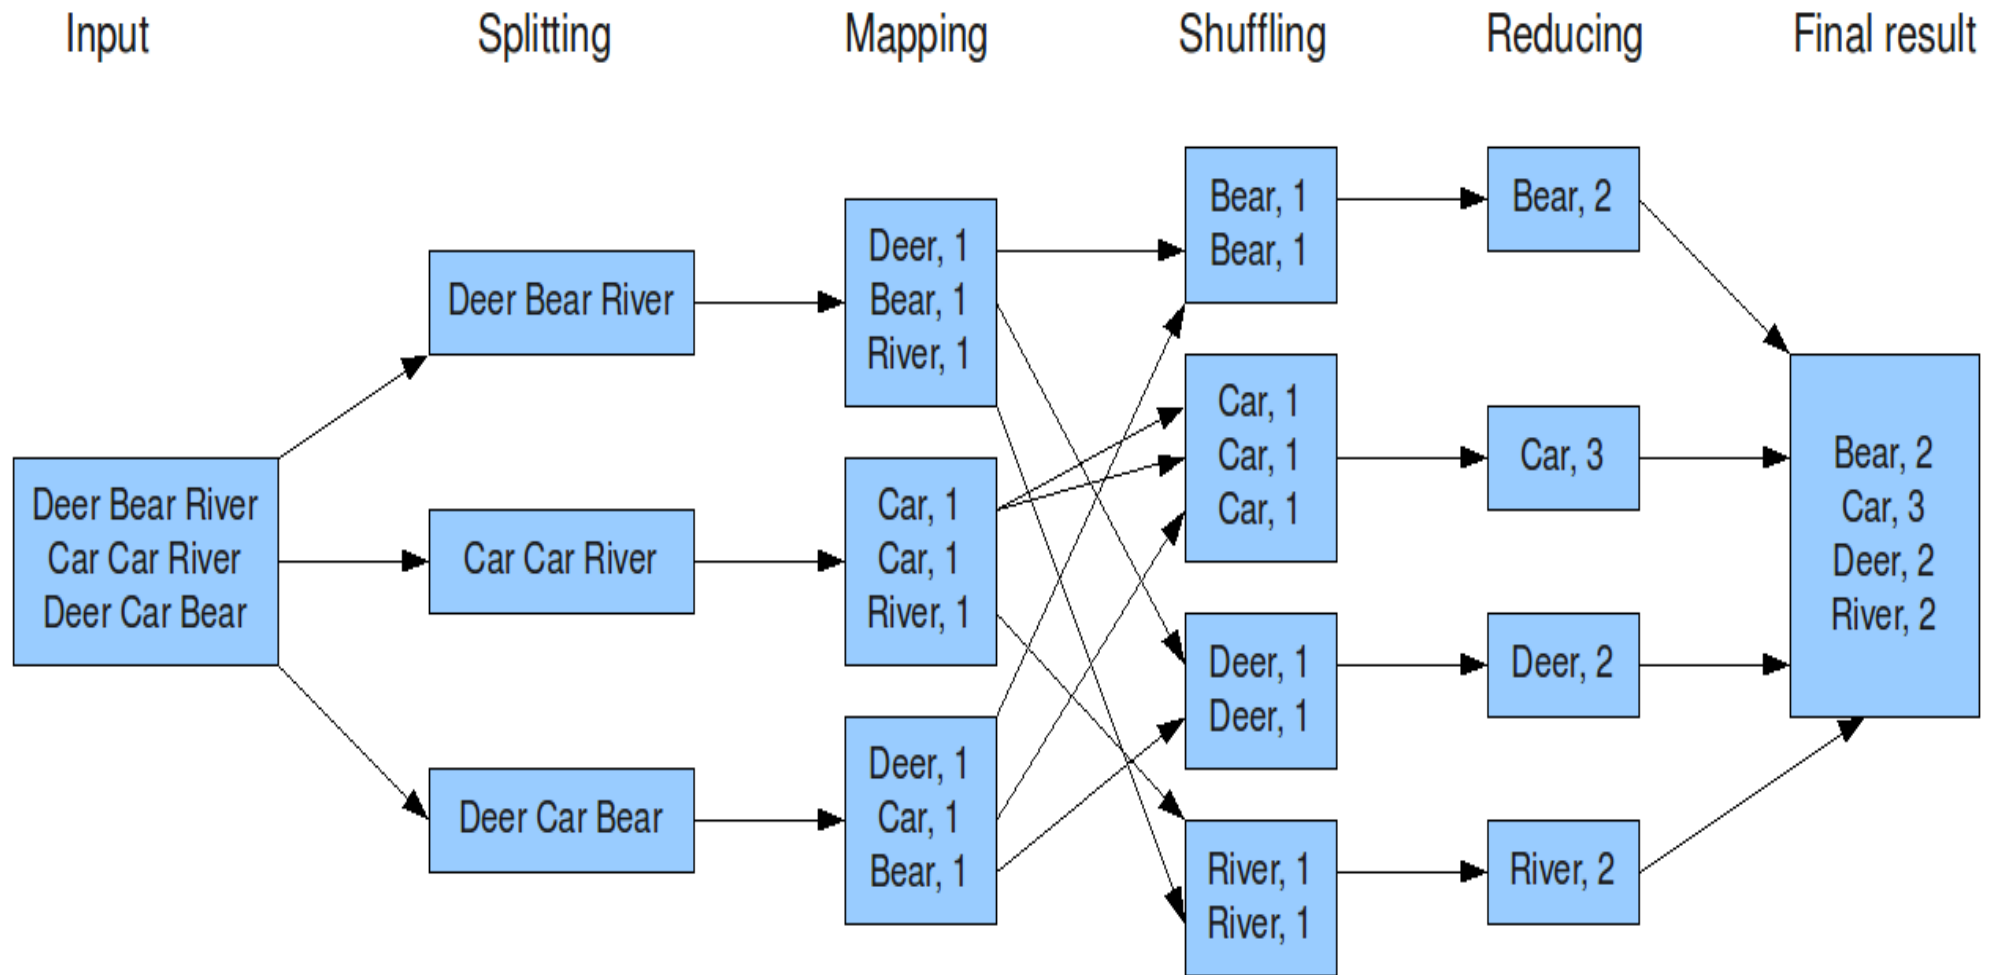
Hadoop Programming with Java

- WordCount
- MaxTemp

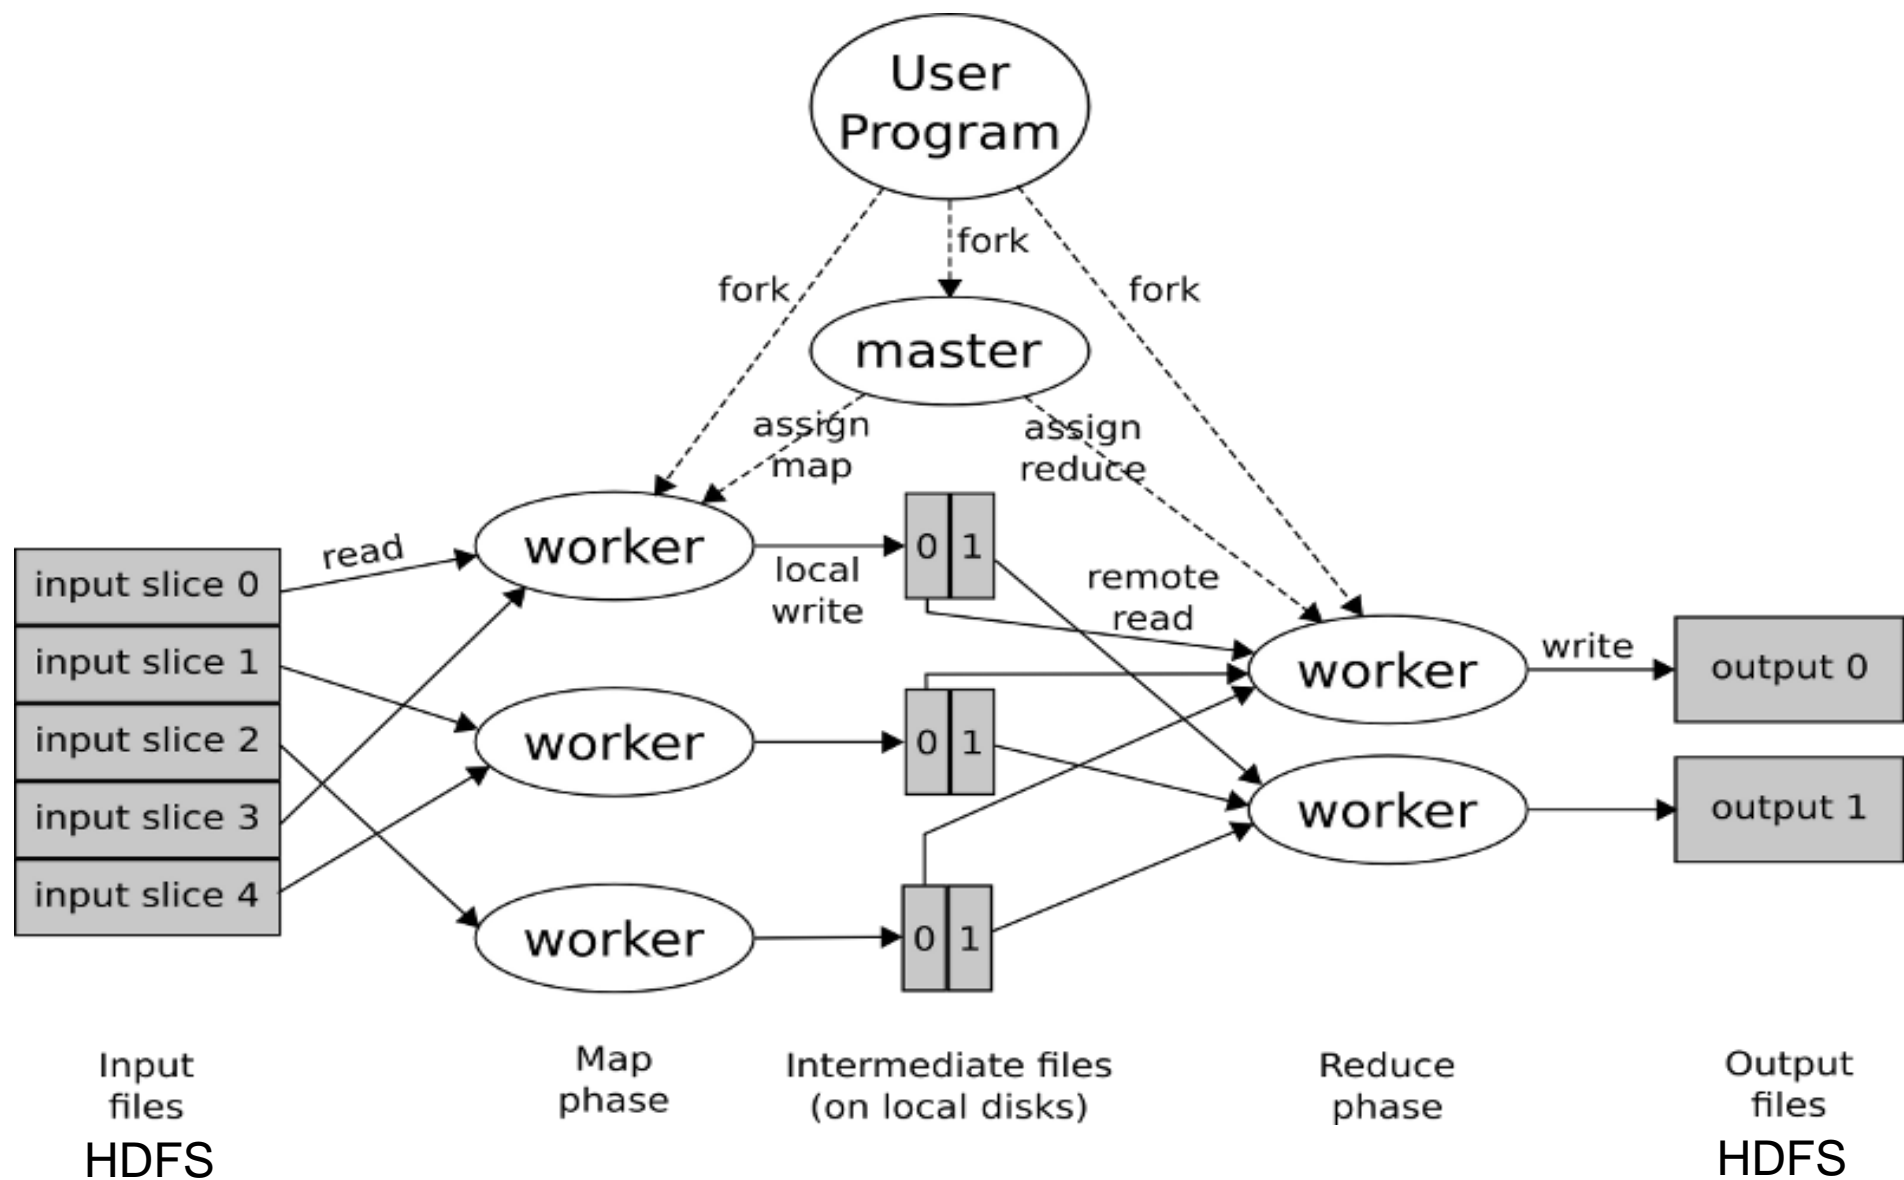
Hadoop Programming with Python

- WordCount
- MaxTemp
- WordLength
- InvertedIndex
- LinkCount

Reminder



Reminder



Target

To be able to write distributed programs over a **Hadoop cluster**.

The examples are simple for illustration purposes BUT the process we will follow is the same either we have an easy or a difficult problem.

Infrastructure

The cluster at LiX is composed of 32 physical nodes, of which one is the **MASTER** and the others are **SLAVES**.

All machines are running **CentOS Linux** and they have 16GB of RAM and 2TB of disk space.

The cluster is used for training purposes.

Prerequisites

To have an Internet connection.

To be able to login to
master-bigdata.polytechnique.fr

To have at least a **small experience** with programming.

Prerequisites

Login to

master-bigdata.polytechnique.fr

using your username/password

Copy the file **/opt/hadooplab.tar.gz** in your home directory by executing:

```
cp /opt/hadooplab.tar.gz ~
```

Extract the archive:

```
tar xvf hadooplab.tar.gz
```


HDFS

To get a list of all available commands

```
hadoop fs -help
```

HDFS

Listing files

```
hadoop fs -ls /
```

```
hadoop fs -ls /user
```

```
hadoop fs -ls /user/username
```

Useful: If you do not specify a path, HDFS will execute the command in your HDFS home directory which is /user/username****

HDFS

Show file contents

```
hadoop fs -cat /dssp/data/leonardo/leonardo.txt
```

HDFS

File copy

```
hadoop fs -cp /dssp/data/leonardo/leonardo.txt /user/username/leonardo.txt
```

OR

```
hadoop fs -cp /dssp/data/leonardo/leonardo.txt leonardo.txt
```

View the file

```
hadoop fs -cat /user/username/leonardo.txt
```

OR

```
hadoop fs -cat leonardo.txt
```

Delete the file

```
hadoop fs -rm /user/username/leonardo.txt
```

OR

```
hadoop fs -rm leonardo.txt
```

HDFS

Creating and deleting directories

```
hadoop fs -mkdir /user/username/d1
```

OR

```
hadoop fs -mkdir d1
```

```
hadoop fs -rmdir /user/username/d1
```

OR

```
hadoop fs -rmdir d1
```

HDFS

Delete a directory and ALL CONTENTS

```
hadoop fs -rm -r /some-directory
```

BE VERY CAREFUL WHEN YOU USE IT!

HDFS

Putting/getting files to/from HDFS

```
hadoop fs -put fname.txt /user/username/input
```

OR

```
hadoop fs -put fname.txt input
```

```
hadoop fs -get /user/username/input/fname.txt
```

OR

```
hadoop fs -get fname.txt
```

HDFS Preparation

Input data

All necessary input data files we are going to use are **already located** in HDFS in the directory:

/dssp/data

List the data directories by executing

hadoop fs -ls /dssp/data

HDFS Preparation

We will create an output directory to store the output of hadoop jobs

```
hadoop fs -mkdir /user/username/output
```

OR

```
hadoop fs -mkdir output
```

Hadoop with Java

We will focus on two examples of Hadoop jobs using the Java programming language.

WordCount: given a collection of text documents, find the number of occurrences of each word in the collection.

MaxTemp: given a file containing temperature measurements, find the maximum temperature recording per year.



WordCount: the mapper

```
public static class TokenizerMapper extends Mapper<Object, Text, Text, IntWritable>{

    private final static IntWritable one = new IntWritable(1);
    private Text word = new Text();

    public void map (Object key, Text value, Context context)
        throws IOException, InterruptedException {
        StringTokenizer itr = new StringTokenizer(value.toString());
        while (itr.hasMoreTokens()) {
            word.set (itr.nextToken());
            context.write (word, one);
        }
    }
}
```

WordCount: the reducer

```
public static class IntSumReducer extends
    Reducer<Text,IntWritable,Text,IntWritable> {

    private IntWritable result = new IntWritable();

    public void reduce(Text key, Iterable<IntWritable> values, Context context)
        throws IOException, InterruptedException {

        int sum = 0;

        for (IntWritable val : values) {
            sum += val.get();
        }

        result.set(sum);

        context.write(key, result);
    }
}
```

WordCount: main function

```
public static void main(String[] args) throws Exception {  
    Configuration conf = new Configuration();  
    Job job = Job.getInstance(conf, "word count");  
    job.setJarByClass(WordCount.class);  
    job.setMapperClass(TokenizerMapper.class);  
    job.setCombinerClass(IntSumReducer.class);  
    job.setReducerClass(IntSumReducer.class);  
    job.setOutputKeyClass(Text.class);  
    job.setOutputValueClass(IntWritable.class);  
    FileInputFormat.addInputPath(job, new Path(args[0]));  
    FileOutputFormat.setOutputPath(job, new Path(args[1]));  
    System.exit(job.waitForCompletion(true) ? 0 : 1);  
}
```

WordCount: compiling the code

Go inside the java-wordcount folder, by executing the following command from your home folder:

```
cd ~/hadooplab/java-wordcount
```

**The relevant code is contained in the file
WordCount.java**

WordCount: compiling the code

To compile the code run the command:

```
javac -classpath "$(yarn classpath)" WordCount.java
```

The file **WordCount.class** must have been produced.

WordCount: building the jar

We will create the file

wc.jar

Please execute

```
jar cf wc.jar WordCount*.class
```

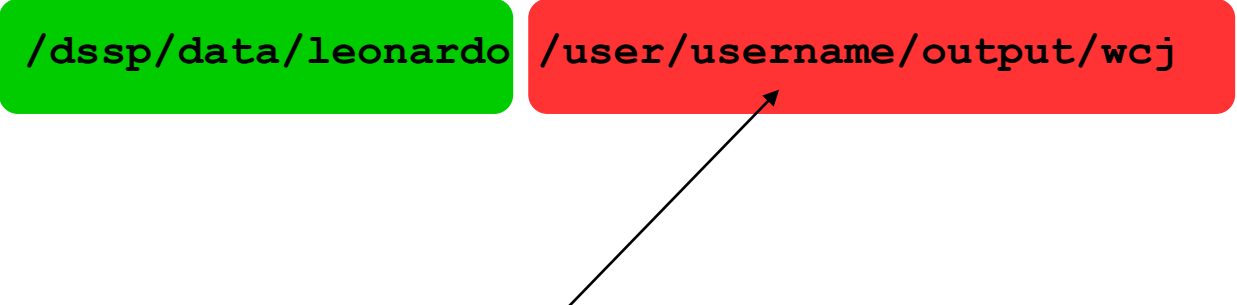
Everything is set! Lets run the job on the cluster.

WordCount: running the job

Execute the following command:

input output

```
hadoop jar wc.jar WordCount /dssp/data/leonardo /user/username/output/wcj
```



Put your **username** here

OR

```
hadoop jar wc.jar WordCount /dssp/data/leonardo output/wcj
```

WordCount: exploring the results

```
hadoop fs -ls output/wcj
```

You should see something like this

```
-rw-r--r-- 3 a.papadopoulos a.papadopoulos      0 2015-05-14 18:32 /user/a.papadopoulos/out2/_SUCCESS
-rw-r--r-- 3 a.papadopoulos a.papadopoulos 53163233 2015-05-14 18:32 /user/a.papadopoulos/out2/part-r-00000
```

WordCount: exploring the results

Examine the last lines of the output:

```
hadoop fs -tail output/wcj/part-r-00000
```

Hadoop with Python

In this part, we focus on the Python language. We will discuss several different problems:

WordCount: given a collection of text documents, find the number of occurrences of each word in the collection.

MaxTemp: given a file containing temperature measurements, find the maximum temperature recording per year.

WordLength: find the average word length in a collection of document beginning by each letter.

InvertedIndex: given a collection of text documents, determine the set of documents that contain each unique word of the collection.

LinkCount: given a graph, determine the number of outgoing links for each node.



Hadoop Streaming

This is a tool that Hadoop provides to support code writing in any language (we will use it for Python)

Assumptions: Mappers should read from **`stdin`** and write to **`stdout`**. Reducers should read from **`stdin`** and write to **`stdout`**.

WordCount: the mapper

```
#!/usr/bin/env python
```

```
import sys
```

```
# input comes from STDIN (standard input)
```

```
for line in sys.stdin:
```

```
    # remove leading and trailing whitespace
```

```
    line = line.strip()
```

```
    # split the line into words
```

```
    words = line.split()
```

```
    # increase counters
```

```
    for word in words:
```

```
        # write the results to STDOUT (standard output);
```

```
        # what we output here will be the input for the
```

```
        # Reduce step, i.e. the input for reducer.py
```

```
        # tab-delimited; the trivial word count is 1
```

```
        print '%s\t%s' % (word, 1)
```

WordCount: running the job

Write the command in a single line!

```
hadoop jar $HADOOP_HOME/share/hadoop/tools/lib/hadoop-streaming-3.1.4.jar -file ./mapper.py -mapper ./mapper.py -file ./reducer.py -reducer ./reducer.py -input /dssp/data/leonardo -output output/wcp
```

WordCount: exploring the results

To view the last lines of the output use **-tail**

```
hadoop fs -tail output/wcp/part-00000
```

To view the whole output file use **-cat**

```
hadoop fs -cat output/wcp/part-00000
```


WordCount: set number of reducers

Write the command in a single line!

```
hadoop jar $HADOOP_HOME/share/hadoop/tools/lib/hadoop-streaming-3.1.4.jar  
-D mapred.reduce.tasks=4  
-file ./mapper.py -mapper ./mapper.py  
-file ./reducer.py -reducer ./reducer.py  
-input /dssp/data/leonardo  
-output output/wcp2
```

WordCount: set number of reducers

Check the number of output files

```
hadoop fs -ls output/wcp2
```

How many output files are there?

WordLength: the problem

Given a document collection find the average word length beginning with each upper case letter.

WordLength: the mapper

```
#!/usr/bin/env python
import sys
for line in sys.stdin:
    line = line.strip()
    words = line.split()
    for word in words:
        if (word[0]>='A') and (word[0] <= 'Z'):
            print '%s\t%s' % (word[0].upper(), len(word))
```

WordLength: running the job

Write the command in a single line!

```
hadoop jar $HADOOP_HOME/share/hadoop/tools/lib/hadoop-streaming-3.1.4.jar  
  -file ./mapper.py -mapper ./mapper.py  
  -file ./reducer.py -reducer ./reducer.py  
  -input /dssp/data/leonardo  
  -output output/wlp
```

WordLength: exploring the results

To view the whole output file use **-cat**

```
hadoop fs -tail output/wlp/part-00000
```

Inversion: the problem

Given a collection of text documents, for each word determine the list of document ids containing the word.

The resulting data structure is known as the **inverted index** and the process of creating it is called **inversion**.

Inversion: running the job

Write the command in a single line!

```
hadoop jar $HADOOP_HOME/share/hadoop/tools/lib/hadoop-streaming-3.1.4.jar  
  -file ./mapper.py -mapper ./mapper.py  
  -file ./reducer.py -reducer ./reducer.py  
  -input /dssp/data/leonardo  
  -output output/invp
```


Inversion: exploring the results

To view the whole output file use **-cat**

```
hadoop fs -tail output/invp/part-00000
```

LinkCount: the problem

Given a graph $G(V,E)$ determine the number of links for each node of the graph.

Web analogy: links are URLs

Input format:

1 4

1 5

2 6

6 8

...

LinkCount: running the job

Write the command in a single line!

```
hadoop jar $HADOOP_HOME/share/hadoop/tools/lib/hadoop-streaming-3.1.4.jar  
  -file ./mapper.py -mapper ./mapper.py  
  -file ./reducer.py -reducer ./reducer.py  
  -input /dssp/data/enron  
  -output output/lcp
```

LinkCount: exploring the results

To view the whole output file use **-cat**

```
hadoop fs -tail output/lcp/part-00000
```

Testing the Code Locally

Since hadoop streaming requires to read from stdin and write to stdout, **we can test our code without submitting the job to the cluster.**

Go into the python-wordcount directory and run the following command:

```
cat ./leonardo.txt | python  
mapper.py | sort | python  
reducer.py > output.txt
```